

Similarity Search

The String Edit Distance

Nikolaus Augsten

nikolaus.augsten@plus.ac.at
Department of Computer Science
University of Salzburg



WS 2022/23

Version January 27, 2023

Outline

1 String Edit Distance

- Motivation and Definition
- Brute Force Algorithm
- Dynamic Programming Algorithm
- Edit Distance Variants

2 Conclusion

Outline

1 String Edit Distance

- Motivation and Definition
- Brute Force Algorithm
- Dynamic Programming Algorithm
- Edit Distance Variants

2 Conclusion

Motivation

- How different are
 - `hello` and `hello?`
 - `hello` and `hallo?`
 - `hello` and `hell?`
 - `hello` and `shell?`

What is a String Distance Function?

Definition (String Distance Function)

Given a finite alphabet Σ , a *string distance function*, δ_s , maps each pair of strings $(x, y) \in \Sigma^* \times \Sigma^*$ to a positive real number (including zero).

$$\delta_s : \Sigma^* \times \Sigma^* \rightarrow \mathbb{R}_0^+$$

- Σ^* is the set of all strings over Σ , including the empty string ε .

The String Edit Distance

Definition (String Edit Distance)

The *string edit distance* between two strings, $\text{ed}(x, y)$, is the minimum number of character insertions, deletions and replacements that transforms x to y .

- Example:
 - $\text{hello} \rightarrow \text{hallo}$: replace e by a
 - $\text{hello} \rightarrow \text{hell}$: delete o
 - $\text{hello} \rightarrow \text{shell}$: delete o , insert s
- Also called *Levenshtein distance*.¹

¹Levenshtein introduced this distance for signal processing in 1965 [Lev65].

Outline

1 String Edit Distance

- Motivation and Definition
- Brute Force Algorithm
- Dynamic Programming Algorithm
- Edit Distance Variants

2 Conclusion

Gap Representation


- **Gap representation** of the string transformation $x \rightarrow y$:
Place string x above string y
 - with a gap in x for every insertion,
 - with a gap in y for every deletion,
 - with different characters in x and y for every replacement.
- Any sequence of edit operations can be represented with gaps.
- **Example:**

h	a	l	l	o
s	h	e	l	l

- insert s
- replace a by e
- delete o

Deriving the Recursive Formula: Optimal Substructure

- **Recursive solution:** is applicable only to problems with optimal substructure property.
- **Optimal substructure** property of a problem:
 - optimal solution to larger problem computable from the optimal solutions of subproblems
- **Examples:**
 - Shortest path has optimal substructure: If a is on shortest path P from a to $b \Rightarrow$ the section $a \rightarrow c$ on P is the shortest path between a and c .
 - Longest simple² path does *not* have optimal substructure. Counter example [CLRS09]: Consider longest path $q \rightarrow t$ and subpath $q \rightarrow r$.



²i.e., the path has no cycles

Deriving the Recursive Formula: Optimal Substructure

Lemma (Optimal Substructure of String Edit Distance Problem)

Given a gap representation, $\text{gap}(x, y)$, between two strings x and y , such that the cost of $\text{gap}(x, y)$ is the string edit distance $\text{ed}(x, y)$.

If we remove the last column of $\text{gap}(x, y)$, then the gap representation of the remaining columns, $\text{gap}(x', y')$, has cost $\text{ed}(x', y')$ between the resulting substrings, x' and y' .

```
h a l l|o  
s h e l l|
```

- **Example:**

- $x = \text{hallo}$, $y = \text{shell}$, $\text{cost}(\text{gap}(x, y)) = \text{ed}(x, y) = 3$
- $x' = \text{hall}$, $y' = \text{shell}$, $\text{cost}(\text{gap}(x', y')) = \text{ed}(x', y') = 2$

Deriving the Recursive Formula: Optimal Substructure

Proof: Optimal Substructure String Edit Distance (by contradiction).

- Last column contributes with $c = 0$ or $c = 1$ to cost of $\text{gap}(x, y)$, thus:

$$\text{cost}(\text{gap}(x, y)) = \text{cost}(\text{gap}(x', y')) + c$$

- Assume $\text{gap}(x', y')$ is not optimal, i.e., $\text{cost}(\text{gap}(x', y')) > \text{ed}(x', y')$. Let $\text{gap}^*(x', y')$ be the respective gap representation:

$$\text{cost}(\text{gap}^*(x', y')) = \text{ed}(x', y') < \text{cost}(\text{gap}(x', y'))$$

- By extending $\text{gap}^*(x', y')$ with the last column, we get a gap representation $\text{gap}^*(x, y)$ with cost below $\text{ed}(x, y)$, which contradicts the definition of the edit distance.

$$\begin{aligned}\text{cost}(\text{gap}^*(x, y)) &= \text{cost}(\text{gap}^*(x', y')) + c \\ &< \text{cost}(\text{gap}(x', y')) + c = \text{ed}(x, y)\end{aligned}$$



Deriving the Recursive Formula

- Example:

```
h a l l o  
s h e l l
```

- Notation:

- $x[1\dots i]$ is the substring of the first i characters of x ($x[1\dots 0] = \varepsilon$)
- $x[i]$ is the i -th character of x

- Recursive Formula:

$$\begin{aligned}\text{ed}(\varepsilon, \varepsilon) &= 0 \\ \text{ed}(x[1..i], \varepsilon) &= i \\ \text{ed}(\varepsilon, y[1..j]) &= j \\ \text{ed}(x[1..i], y[1..j]) &= \min(\text{ed}(x[1..i-1], y[1..j-1]) + c, \\ &\quad \text{ed}(x[1..i-1], y[1..j]) + 1, \\ &\quad \text{ed}(x[1..i], y[1..j-1]) + 1)\end{aligned}$$

where $c = 0$ if $x[i] = y[j]$, otherwise $c = 1$.


Brute Force Algorithm

$\text{ed-bf}(x, y)$

```
 $m = |x|, n = |y|$ 
if  $m = 0$  then return  $n$ 
if  $n = 0$  then return  $m$ 
if  $x[m] = y[n]$  then  $c = 0$  else  $c = 1$ 
return  $\min(\text{ed-bf}(x, y[1 \dots n - 1]) + 1,$ 
         $\text{ed-bf}(x[1 \dots m - 1], y) + 1,$ 
         $\text{ed-bf}(x[1 \dots m - 1], y[1 \dots n - 1]) + c)$ 
```

Brute Force Algorithm

- Recursion tree for $\text{ed-bf}(\text{ab}, \text{xb})$:



- Exponential runtime in string length :-(
- **Observation:** Subproblems are computed repeatedly
(e.g. $\text{ed-bf}(\text{a}, \text{x})$ is computed 3 times)
- **Approach:** Reuse previously computed results!

Outline

1 String Edit Distance

- Motivation and Definition
- Brute Force Algorithm
- **Dynamic Programming Algorithm**
- Edit Distance Variants

2 Conclusion

Dynamic Programming Algorithm – Top Down


- Store distances between all prefixes of x and y
- Use matrix $C_{0..m,0..n}$ with

$$C_{i,j} = \text{ed}(x[1 \dots i], y[1 \dots j])$$

where $x[1..0] = y[1..0] = \varepsilon$.

- Example:

	ε	x	b
ε	0	1	2
a	1	1	2
b	2	2	1



Dynamic Programming Algorithm – Bottom Up

ed-dyn(x, y)

```
C : array[0..|x|][0..|y|]
for i = 0 to |x| do C[i, 0] = i
for j = 1 to |y| do C[0, j] = j
for j = 1 to |y| do
    for i = 1 to |x| do
        if x[i] = y[j] then c = 0 else c = 1
        C[i, j] = min(C[i - 1, j - 1] + c,
                        C[i - 1, j] + 1,
                        C[i, j - 1] + 1)
```

Understanding the Solution

- Example:

$x = \text{moon}$
 $y = \text{mond}$

		ins →					
		ren ↓	ε	m	o	n	d
del	↓	ε	0 ← 1 ← 2 ← 3 ← 4				
		m	1 ↑ 0 ← 1 ← 2 ← 3				
o	↓	2 ↑ 1	0 ← 1 ← 2				
		o	3 ↑ 2 ↑ 1	0 ← 1 ← 2			
n	↓	4 ↑ 3	2 ↑ 1	1	0 ← 1 ← 2		
		n			2 ↑ 1	0 ← 1 ← 2	

- Each arrow represents an edit operation with minimal cost
- Cost 2 in cell (n, d) can either result from replacing n by d (diagonal arrow) or by inserting d (horizontal arrow)
- Each path from bottom right to top left corner represents a valid set of edit operations

Understanding the Solution

- Example:

		ins →					m o o n	
							m o n d	
		ε	m	o	n	d		
x =	moon	ε	0	1	2	3	4	
y =	mond	del ↓ m	1	0	1	2	3	m o o n
		o	2	1	0	1	2	m o n d
		o	3	2	1	1	2	m o o n
		n	4	3	2	1	2	m o n d

- Solution 1: replace **n** by **d** and (second) **o** by **n** in **x**
- Solution 2: insert **d** after **n** and delete (first) **o** in **x**
- Solution 3: insert **d** after **n** and delete (second) **o** in **x**

Dynamic Programming Algorithm

ed-dyn⁺(x, y)

```
col0 : array[0..|x|]
col1 : array[0..|x|]
for i = 0 to |x| do col0[i] = i
for j = 1 to |y| do
    col1[0] = j
    for i = 1 to |x| do
        if x[i] = y[j] then c = 0 else c = 1
        col1[i] = min(col0[i - 1] + c,
                        col1[i - 1] + 1,
                        col0[i] + 1)
    col0 = col1
```

Outline

1 String Edit Distance

- Motivation and Definition
- Brute Force Algorithm
- Dynamic Programming Algorithm
- Edit Distance Variants

2 Conclusion

Distance Metric

Definition (Distance Metric)

A distance function δ is a *distance metric* if and only if for any x, y, z the following hold:

- $\delta(x, y) = 0 \Leftrightarrow x = y$ (identity)
- $\delta(x, y) = \delta(y, x)$ (symmetric)
- $\delta(x, y) + \delta(y, z) \geq \delta(x, z)$ (triangle inequality)

Examples:

- the Euclidean distance is a metric
- $d(a, b) = a - b$ is not a metric (not symmetric)

Introducing Weights

- Look at the edit operations as a set of rules with a cost:

$$\begin{aligned}\alpha(\varepsilon, b) &= \omega_{ins} && (\text{insert}) \\ \alpha(a, \varepsilon) &= \omega_{del} && (\text{delete}) \\ \alpha(a, b) &= \begin{cases} \omega_{rep} & \text{if } a \neq b \\ 0 & \text{if } a = b \end{cases} && (\text{replace})\end{aligned}$$

where $a, b \in \Sigma$, and $\omega_{ins}, \omega_{del}, \omega_{rep} \in \mathbb{R}_0^+$.

- **Edit script**: sequence of rules that transform x to y
- **Edit distance**: edit script with minimum cost
(adding up costs of single rules)
- **Example**: so far we assumed $\omega_{ins} = \omega_{del} = \omega_{rep} = 1$.

Weighted Edit Distance

- Recursive formula with weights:

$$C_{0,0} = 0$$

$$\begin{aligned} C_{i,j} = \min(& C_{i-1,j-1} + \alpha(x[i], y[j]), \\ & C_{i-1,j} + \alpha(x[i], \varepsilon), \\ & C_{i,j-1} + \alpha(\varepsilon, y[j])) \end{aligned}$$

where $\alpha(a, a) = 0$ for all $a \in \Sigma$, and $C_{-1,j} = C_{i,-1} = \infty$.

- We can easily adapt the dynamic programming algorithm.

Variants of the Edit Distance

- Unit cost edit distance (what we did so far):
 - $\omega_{ins} = \omega_{del} = \omega_{rep} = 1$
 - $0 \leq ed(x, y) \leq \max(|x|, |y|)$
 - distance metric
- Hamming distance [Ham50, SK83]:
 - called also “string matching with k mismatches”
 - allows only replacements
 - $\omega_{rep} = 1, \omega_{ins} = \omega_{del} = \infty$
 - $0 \leq d(x, y) \leq |x|$ if $|x| = |y|$, otherwise $d(x, y) = \infty$
 - distance metric
- Longest Common Subsequence (LCS) distance [NW70, AG87]:
 - allows only insertions and deletions
 - $\omega_{ins} = \omega_{del} = 1, \omega_{rep} = \infty$
 - $0 \leq d(x, y) \leq |x| + |y|$
 - distance metric
 - $LCS(x, y) = (|x| + |y| - d(x, y))/2$

Allowing Transposition

- Transpositions
 - switch two adjacent characters
 - can be simulated by delete and insert
 - typos are often transpositions
- New rule for transposition

$$\alpha(ab, ba) = \omega_{trans}$$

allows us to assign a weight different from $\omega_{ins} + \omega_{del}$

- Recursive formula that includes transposition:

$$C_{0,0} = 0$$

$$\begin{aligned} C_{i,j} = \min(& C_{i-1,j-1} + \alpha(x[i], y[j]), \\ & C_{i-1,j} + \alpha(x[i], \varepsilon), \\ & C_{i,j-1} + \alpha(\varepsilon, y[j]), \\ & C_{i-2,j-2} + \alpha(x[i-1]x[i], y[j-1]y[j])) \end{aligned}$$

where $\alpha(ab, cd) = \infty$ if $a \neq d$ or $b \neq c$, $\alpha(a, a) = 0$ for all $a \in \Sigma$,
and $C_{-1,j} = C_{i,-1} = C_{-2,j} = C_{i,-2} = \infty$.

Example: Edit Distance with Transposition

- Example: Compute distance between $x = \text{meal}$ and $y = \text{mael}$ using the edit distance with transposition ($\omega_{ins} = \omega_{del} = \omega_{rep} = \omega_{trans} = 1$)

	ε	m	a	e	l
ε	0	1	2	3	4
m	1	0	1	2	3
e	2	1	1	1	2
a	3	2	1	1	2
l	4	3	2	2	1

- The value in red results from the transposition of ea to ae.

Text Searching

- Goal:
 - search pattern p in text t ($|p| < |t|$)
 - allow k errors
 - match may start at any position of the text
- Difference to distance computation:
 - $C_{0,j} = 0$ (instead of $C_{0,j} = j$, as text may start at any position)
 - result: all $C_{m,j} \leq k$ are endpoints of matches

Example: Text Searching

- Example:

$p = \text{survey}$
 $t = \text{surgery}$
 $k = 2$

	ε	s	u	r	g	e	r	y
ε	0	0	0	0	0	0	0	0
s	1	0	1	1	1	1	1	1
u	2	1	0	1	2	2	2	2
r	3	2	1	0	1	2	2	3
v	4	3	2	1	1	2	3	3
e	5	4	3	2	2	1	2	3
y	6	5	4	3	3	2	2	2

- Solutions: 3 matching positions with $k \leq 2$ found.

s u r v e y
s u r g e

s u r v e y
s u r g e r

s u r v e y
s u r g e r y

Summary

- Edit distance between two strings: the minimum number of edit operations that transforms one string into the another
- Dynamic programming algorithm with $O(mn)$ time and $O(m)$ space complexity, where $m \leq n$ are the string lengths.
- Basic algorithm can easily be extended in order to:
 - weight edit operations differently,
 - support transposition,
 - simulate Hamming distance and LCS,
 - search pattern in text with k errors.

-  Alberto Apostolico and Zvi Galil.
The longest common subsequence problem revisited.
Algorithmica, 2(1):315–336, March 1987.
-  Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein.
Introduction to Algorithms, 3rd Edition.
MIT Press, 2009.
-  Richard W. Hamming.
Error detecting and error correcting codes.
Bell System Technical Journal, 26(2):147–160, 1950.
-  Vladimir I. Levenshtein.
Binary codes capable of correcting spurious insertions and deletions of ones.
Problems of Information Transmission, 1:8–17, 1965.
-  Saul B. Needleman and Christian D. Wunsch.

A general method applicable to the search for similarities in the amino acid sequence of two proteins.

Journal of Molecular Biology, 48:443–453, 1970.



David Sankoff and Josef B. Kruskal, editors.

Time Warps, String Edits, and Macromolecules: The Theory and Practice of Sequence Comparison.

Addison-Wesley, Reading, MA, 1983.